

FactoryElectricPrediction

April 15, 2025

1 Factory Electric Consumption Prediction

Forecast the electric consumption of a factory based on historical data. The goal is to develop a predictive model that provides accurate electricity consumption forecasts to improve energy efficiency, plan resource allocation, and identify seasonal patterns or trends.
Objective - Predict Factory Electric Consumption using a regression model.
- Training data: 13,872 records - Test data: 2,160 records - Evaluation metric: RMSE

1.1 Libraries Used

```
[1]: import pandas as pd
import numpy as np
import os
from xgboost import XGBRegressor
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.linear_model import LinearRegression
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import PolynomialFeatures
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
import seaborn as sns
```

1.2 Data Exploration

```
[2]: # Upload the data
train_df = pd.read_csv("../Data/train_df.csv")
test_df = pd.read_csv("../Data/test_df.csv")


# Convert Date column to datetime
train_df['Date'] = pd.to_datetime(train_df['Date'])
test_df['Date'] = pd.to_datetime(test_df['Date'])

# Statistics
display(train_df.describe())
```

	Electric_Consumption	Factor_A	Factor_B	Factor_C \
count	13872.000000	13872.000000	13872.000000	13872.000000
mean	5.893120	2.262227	19.977867	40.332796
std	7.355969	1.457234	9.607866	24.324671
min	-0.009717	-74.220598	-5.422289	3.249366
25%	0.000000	1.105756	13.088434	20.396742
50%	0.086858	1.995936	20.187978	34.024351
75%	13.512742	3.261530	26.994172	57.191932
max	21.360638	8.392045	44.036205	101.706771

	Factor_D	Factor_E	Factor_F
count	13872.000000	13872.000000	13872.000000
mean	145.005668	0.209435	42.522214
std	94.993992	1.691576	280.466514
min	-1380.363752	0.000000	0.000000
25%	106.883078	0.000000	0.579665
50%	138.044038	0.000000	1.981719
75%	209.901463	0.000000	60.123072
max	312.895406	31.000006	31839.840610

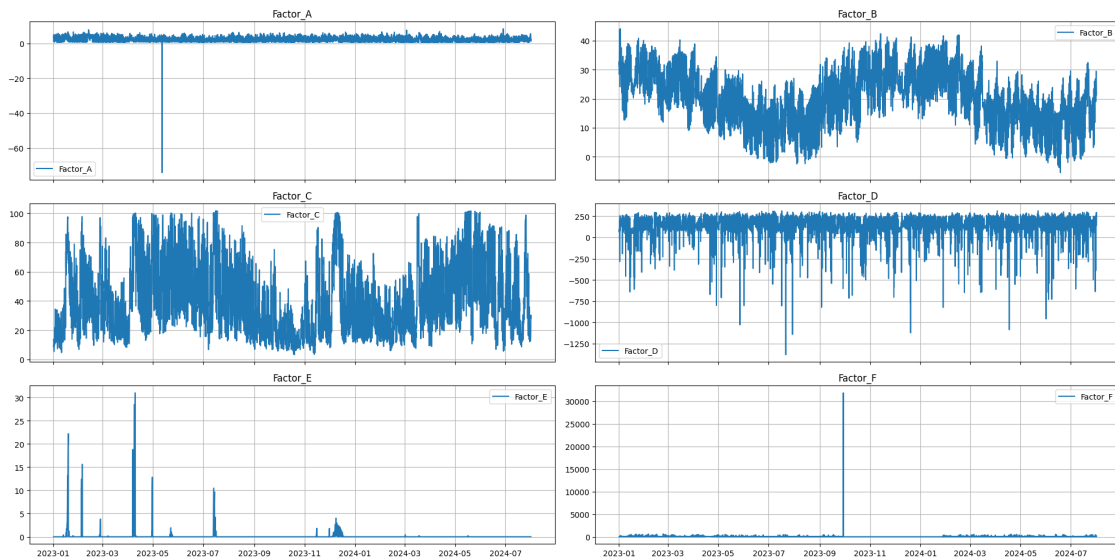
```
[3]: # Plot the Electric Consumption over time
plt.figure(figsize=(10, 5))
plt.plot(train_df['Date'], train_df['Electric_Consumption'], label='Electric_
↳Consumption')
plt.title('Electric Consumption Over Time')
plt.xlabel('Date')
plt.ylabel('Electric Consumption')
plt.legend()
plt.show()
```



```
[4]: # Feature Visualization: create subplots for different factors over time
fig, axs = plt.subplots(3, 2, figsize=(20, 10), sharex=True)
axs = axs.flatten()

factors = ['Factor_A', 'Factor_B', 'Factor_C', 'Factor_D', 'Factor_E', 'Factor_F']

for i, factor in enumerate(factors):
    axs[i].plot(train_df['Date'], train_df[factor], label=factor)
    axs[i].set_title(factor)
    axs[i].legend()
    axs[i].grid(True)
plt.tight_layout()
plt.show()
```



```
[5]: # Feature Extraction
for df in [train_df, test_df]:
    df['Hour'] = df['Date'].dt.hour
    df['DayOfWeek'] = df['Date'].dt.dayofweek
    df['Month'] = df['Date'].dt.month
    df['Day'] = df['Date'].dt.day

test_dates = pd.to_datetime(test_df['Date'])

# Drop the original Date column
train_df.drop('Date', axis=1, inplace=True)
```

```
test_df.drop('Date', axis=1, inplace=True)

features = ['Factor_A', 'Factor_B', 'Factor_C', 'Factor_D', 'Factor_E',
            'Factor_F', 'Hour', 'DayOfWeek', 'Month', 'Day']

display(train_df.head())
```

	Electric_Consumption	Factor_A	Factor_B	Factor_C	Factor_D	Factor_E	\
0	0.000000	1.242130	28.419739	13.720397	79.840600	0.0	
1	0.000000	1.861285	29.840759	12.537668	86.424903	0.0	
2	0.000000	4.212674	32.778036	9.408667	72.082793	0.0	
3	0.000000	4.025251	32.624700	9.035601	73.825705	0.0	
4	-0.000267	3.122659	31.931245	9.235502	66.823956	0.0	

	Factor_F	Hour	DayOfWeek	Month	Day
0	2.386157	0	6	1	1
1	1.473256	1	6	1	1
2	1.583711	2	6	1	1
3	1.706656	3	6	1	1
4	0.987048	4	6	1	1

```
[6]: # Missing Values Check
train_df.isnull().sum().sum()
```

[6]: 0

1.3 Outliers Analysis

```
[7]: # Helper functions to display outliers
def boxplt_display(factors):
    fig, axes = plt.subplots(2, 3, figsize=(18, 10))
    fig.suptitle('Boxplots for Each Factor', fontsize=18)
    axes = axes.flatten()

    for i, col in enumerate(factors):
        sns.boxplot(y=train_df[col], ax=axes[i], color='skyblue')
        axes[i].set_title(col, fontsize=12)
        axes[i].set_xlabel('')
        axes[i].set_ylabel('Value')

    plt.tight_layout(rect=[0, 0.03, 1, 0.95])
    plt.show()

def distribution_display(factors):
    fig, axes = plt.subplots(2, 3, figsize=(18, 10))
    fig.suptitle('Distribution Plots for Each Factor', fontsize=18)
    axes = axes.flatten()
```

```

for i, col in enumerate(factors):
    sns.histplot(train_df[col], kde=True, ax=axes[i], color='steelblue')
    axes[i].set_title(f'{col}', fontsize=12)
    axes[i].set_xlabel('')
    axes[i].set_ylabel('Frequency')

for j in range(len(factors), len(axes)):
    fig.delaxes(axes[j])

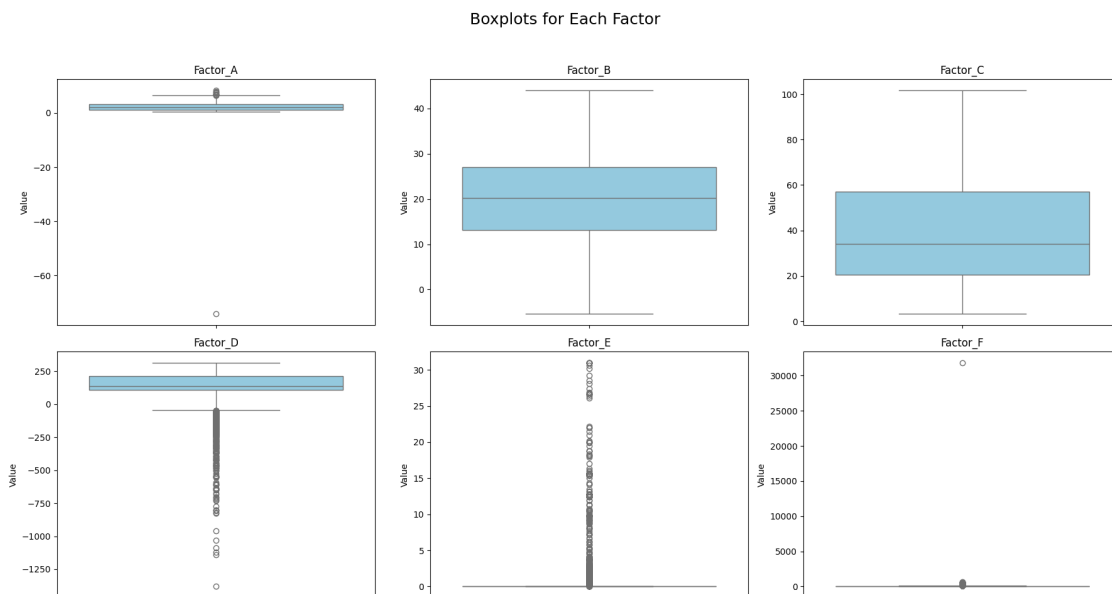
plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.show()

```

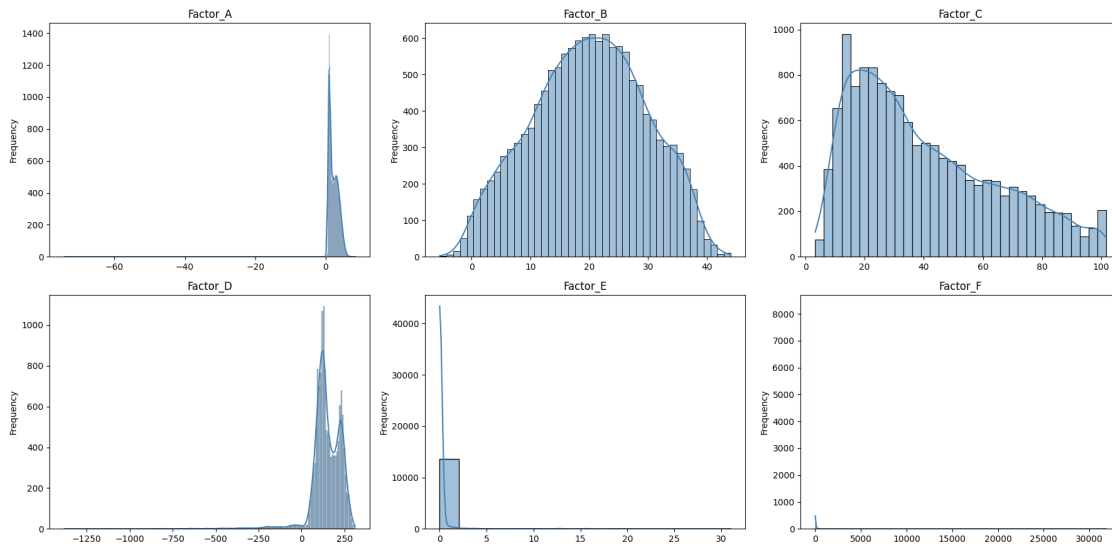
```

[8]: boxplt_display(factors)
distribution_display(factors)

```



Distribution Plots for Each Factor



```
[9]: # Detection and Treatment of Outliers in some Factors through IQR method
# Assuming that Factor_A, Factor_D and Factor_F are the ones with outliers
    ↳ based on previous boxplot and distribution plots analysis

for col in ['Factor_A', 'Factor_D', 'Factor_F']:
    # IQR calculation
    Q1 = train_df[col].quantile(0.25)
    Q3 = train_df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR

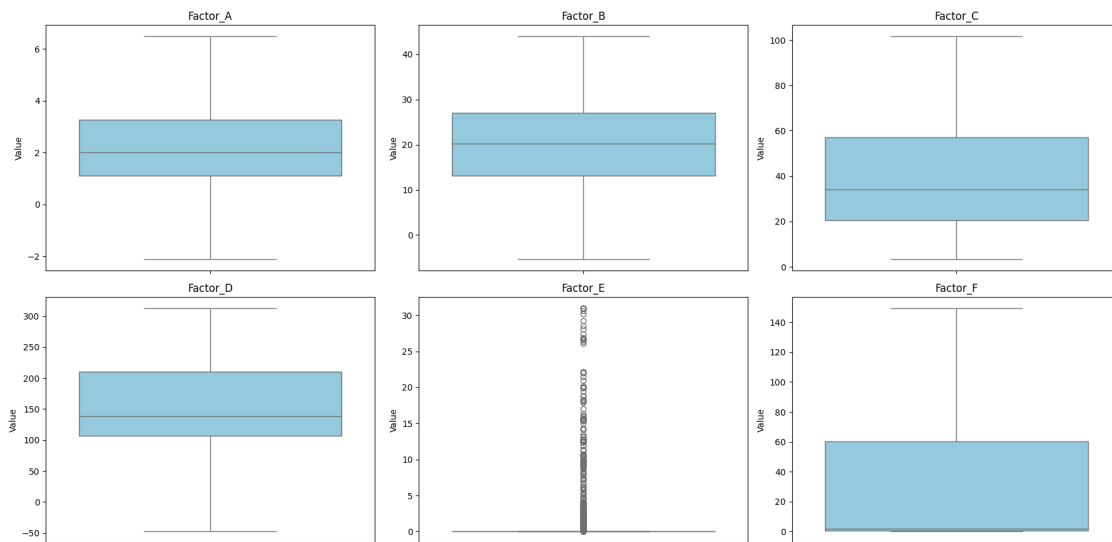
    # Option 1: Clip outliers
    train_df[col] = train_df[col].clip(lower=lower, upper=upper)

    # Option 2: Apply log transform.
    # train_df[col] = np.log1p(train_df[col] - train_df[col].min() + 1)

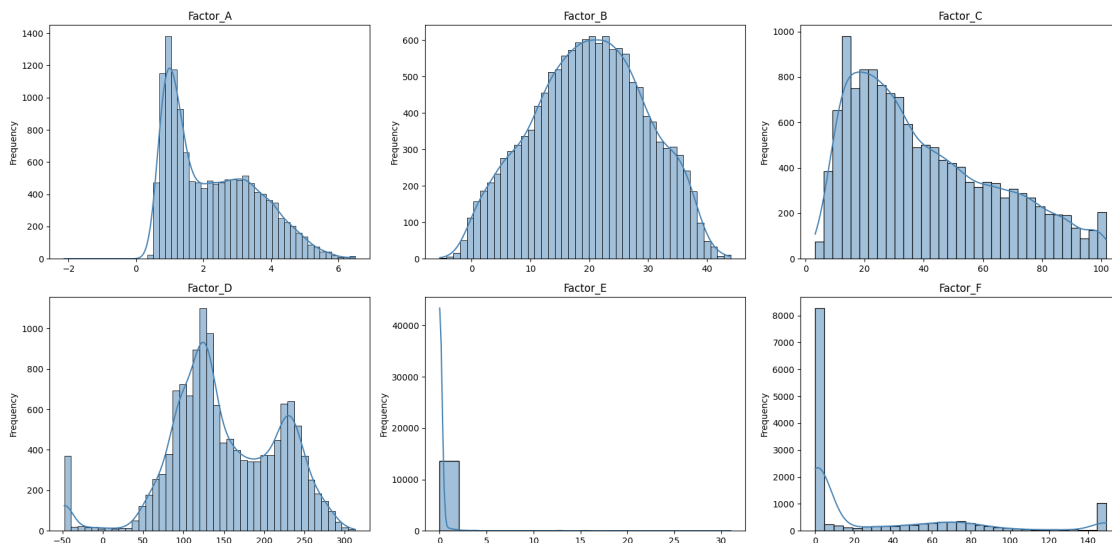
    # Trying both methods to see which one works better but the second option
    ↳ raised slightly the RMSE

# Check for outliers again
boxplt_display(factors)
distribution_display(factors)
```

Boxplots for Each Factor



Distribution Plots for Each Factor



1.4 Models Training and Evaluation

```
[10]: # Dataset Preparation
X = train_df[features]
y = train_df['Electric_Consumption']
X_test = test_df
```

```
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2,
↪random_state=42)

X_test.head()
```

```
[10]:
```

	Factor_A	Factor_B	Factor_C	Factor_D	Factor_E	Factor_F	Hour	\
0	1.775026	21.729808	24.808146	249.474701	0.0	1.808403	0	
1	2.176429	20.792287	25.128845	241.233210	0.0	1.847753	1	
2	2.644089	20.041586	25.045506	239.540034	0.0	1.967446	2	
3	2.759897	18.551710	26.976024	238.425577	0.0	2.128126	3	
4	2.670419	16.689420	29.611734	240.139421	0.0	1.945275	4	

	DayOfWeek	Month	Day
0	3	8	1
1	3	8	1
2	3	8	1
3	3	8	1
4	3	8	1

```
[11]: # Helper Functions
def rmse_cv(model, X, y):
    rmse = np.sqrt(-cross_val_score(model, X, y,
↪scoring="neg_mean_squared_error", cv=5))
    return rmse

def plot_feature_importance(model, features, model_name):
    importances = model.feature_importances_
    indices = np.argsort(importances)[::-1]

    plt.figure(figsize=(10, 6))
    plt.title(f"Feature Importances: {model_name}")
    plt.bar(range(len(importances)), importances[indices])
    plt.xticks(range(len(importances)), [features[i] for i in indices],
↪rotation=45)
    plt.tight_layout()
    plt.show()

def submission_dataset(predictions, name):
    name = name.replace(" ", "_")
    output_file = f"..\\Data\\submission_data\\submission_{name}.csv"
    if os.path.exists(output_file):
        os.remove(output_file)
    submission = pd.DataFrame({
        'Date': test_dates,
        'Electric_Consumption': predictions
    })
    submission.to_csv(output_file, index=False)
```



```
# print(f"File submission_{name}.csv salvato con successo!")
```

```
[12]: # Models to Compare
models = {
    "Linear Regression": LinearRegression(),
    "Polynomial Regression": make_pipeline(PolynomialFeatures(),
    ↪LinearRegression()),
    "Random Forest": RandomForestRegressor(),
    "XGBoost": XGBRegressor(),
}

pr_param_grid = {
    'polynomialfeatures__degree': [2, 3],
    'polynomialfeatures__include_bias': [False, True],
}

rf_param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [6, 10, None],
    'min_samples_split': [2, 5],
}

xgb_param_grid = {
    'n_estimators': [500, 1000],
    'learning_rate': [0.01, 0.05],
    'max_depth': [4, 6, 8],
    'subsample': [0.8, 1],
    'tree_method': ['hist'],
    'predictor': ['cpu_predictor'],
}
```

```
[13]: display(X_train.head())
display(X_test.head())
print("X_train shape",X_train.shape)
print("X_test shape",X_test.shape)
print("y shape",y.shape)
```

	Factor_A	Factor_B	Factor_C	Factor_D	Factor_E	Factor_F	\
9205	3.358914	36.536203	12.820157	145.295411	0.0	0.000000e+00	
6779	2.642509	34.672805	6.255283	251.108846	0.0	2.090000e-43	
4148	1.239516	14.581460	41.370703	129.642411	0.0	1.550142e+00	
4436	0.724160	7.163648	36.193505	233.031004	0.0	1.812564e+00	
13219	1.676083	11.344591	39.577381	136.578953	0.0	1.976420e+00	

	Hour	DayOfWeek	Month	Day
9205	13	4	1	19
6779	11	1	10	10

4148	20	3	6	22
4436	20	1	7	4
13219	19	3	7	4

	Factor_A	Factor_B	Factor_C	Factor_D	Factor_E	Factor_F	Hour	\
0	1.775026	21.729808	24.808146	249.474701	0.0	1.808403	0	
1	2.176429	20.792287	25.128845	241.233210	0.0	1.847753	1	
2	2.644089	20.041586	25.045506	239.540034	0.0	1.967446	2	
3	2.759897	18.551710	26.976024	238.425577	0.0	2.128126	3	
4	2.670419	16.689420	29.611734	240.139421	0.0	1.945275	4	

	DayOfWeek	Month	Day
0	3	8	1
1	3	8	1
2	3	8	1
3	3	8	1
4	3	8	1

X_train shape (11097, 10)
X_test shape (2160, 10)
y shape (13872,)

```
[14]: # Training, Validation, and Results
results_dict = {}

for name, model in models.items():

    print(f"\n##### Training {name} #####\n")
    results_dict[name] = {}

    # Cross validation RMSE
    cv_rmse = rmse_cv(model, X, y)
    print(f"{name} Cross-validated RMSE: {cv_rmse.mean():.2f} ± {cv_rmse.std():.2f}")
    results_dict[name]["cv_rmse"] = cv_rmse.mean()
    results_dict[name]["cv_rmse_std"] = cv_rmse.std()

    if name == "Random Forest": param_grid = rf_param_grid
    elif name == "XGBoost": param_grid = xgb_param_grid
    elif name == "Polynomial Regression": param_grid = pr_param_grid
    else: param_grid = None

    if param_grid:

        # Hyperparameter tuning using GridSearchCV
        print(f"Performing grid search for {name} ...")
        grid_model = GridSearchCV(model, param_grid, cv=3,
        scoring='neg_mean_squared_error', n_jobs=-1, verbose=1, error_score='raise')
```

```

    # Fit the model
    grid_model.fit(X_train, y_train)
    best_model = grid_model.best_estimator_
    best_params = grid_model.best_params_
    print(f"Best {name} parameters: {best_params}")

    results_dict[name]["best_params"] = best_params

else:
    # Fit the model without hyperparameter tuning
    model.fit(X_train, y_train)
    best_model = model
    results_dict[name]["best_params"] = "Default"

val_predictions = best_model.predict(X_val)
val_predictions = np.clip(val_predictions, 0, None) # Ensure predictions
↳are non-negative (We actually have no explanations about the dataset so we
↳can't be sure that negative values are not possible)

rmse = np.sqrt(mean_squared_error(y_val, val_predictions))
print(f"{name} RMSE: {rmse}")
results_dict[name]["rmse"] = rmse

# Feature Importance
if hasattr(best_model, 'feature_importances_'):
    plot_feature_importance(best_model, features, name)

    # I attempted to improve model performance by removing less important
    ↳features based on feature importance scores.
    # Specifically, I experimented with thresholds of 0.05, 0.02, and 0.01 to
    ↳filter out features with low importance values.
    # After selecting the most relevant features using each threshold, I
    ↳retrained the model and evaluated the new RMSE.
    # However, in all cases, the RMSE increased compared to the original model
    ↳using all features.
    # As a result, I decided to keep the full feature set for training, as
    ↳removing the lower-importance features consistently led to worse performance.

    # important_features = [feature for feature, importance in
    ↳zip(features, best_model.feature_importances_)
    #     if importance > 0.01]

    # if len(important_features) > 0:
    #     # New dataset with important features
    #     X_train_imp = X_train[important_features]

```

```

#         X_val_imp = X_val[important_features]
#         X_test_imp = X_test[important_features]

#         # Re-train the model with important features
#         best_model.fit(X_train_imp, y_train)
#         val_predictions = best_model.predict(X_val_imp)
#         print(f"RMSE with selected features: {np.
↪sqrt(mean_squared_error(y_val, val_predictions)).2f}")

#         predictions = best_model.predict(X_test_imp)
#     else:
#         predictions = best_model.predict(X_test)
# else:

predictions = best_model.predict(X_test)

predictions = np.clip(predictions, 0, None)
submission_dataset(predictions, name)

```

Training Linear Regression

Linear Regression Cross-validated RMSE: 4.86 ± 0.66

Linear Regression RMSE: 4.509047174785689

Training Polynomial Regression

Polynomial Regression Cross-validated RMSE: 4.19 ± 1.05

Performing grid search for Polynomial Regression ...

Fitting 3 folds for each of 4 candidates, totalling 12 fits

Best Polynomial Regression parameters: {'polynomialfeatures__degree': 3,
'polynomialfeatures__include_bias': False}

Polynomial Regression RMSE: 2.3970789723234143

Training Random Forest


Random Forest Cross-validated RMSE: 2.58 ± 0.54

Performing grid search for Random Forest ...

Fitting 3 folds for each of 12 candidates, totalling 36 fits

Best Random Forest parameters: {'max_depth': None, 'min_samples_split': 2,
'n_estimators': 200}

Random Forest RMSE: 1.633206614840841



Training XGBoost

XGBoost Cross-validated RMSE: 2.78 ± 0.56


Performing grid search for XGBoost ...

Fitting 3 folds for each of 24 candidates, totalling 72 fits

c:\Users\irebu\AppData\Local\Programs\Python\Python39\lib\site-packages\xgboost\core.py:158: UserWarning: [16:07:56] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-08cbc0333d8d4aae1-1\xgboost\xgboost-ci-windows\src\learner.cc:740: Parameters: { "predictor" } are not used.

warnings.warn(smsg, UserWarning)

Best XGBoost parameters: {'learning_rate': 0.01, 'max_depth': 8, 'n_estimators': 1000, 'predictor': 'cpu_predictor', 'subsample': 0.8, 'tree_method': 'hist'}
XGBoost RMSE: 1.5239462594591227



```
[15]: # Get feature importances on the best optimized model - XGBoost
feature_importance = best_model.feature_importances_
features = X.columns

# Create a DataFrame with feature importances of the best model
importance_df = pd.DataFrame({'Feature': features, 'Importance': feature_importance})
importance_df = importance_df.sort_values(by='Importance', ascending=False)


display(importance_df)

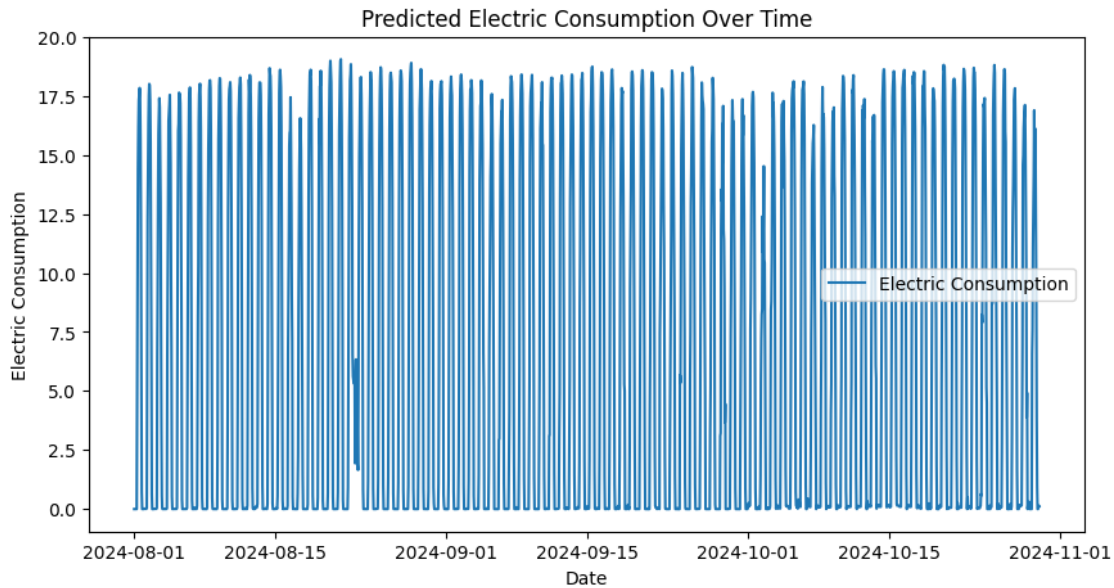
# Plot the features by importance using Seaborn of the best model
plt.figure(figsize=(12, 6))
sns.barplot(x='Importance', y='Feature', data=importance_df.head(5))
plt.title('Top 5 Feature Importances')
plt.xlabel('Importance Score')
plt.ylabel('Features')
plt.tight_layout()
plt.show()

# Plot the predictions
plt.figure(figsize=(10, 5))
plt.plot(test_dates, predictions, label='Electric Consumption')
plt.title('Predicted Electric Consumption Over Time')
plt.xlabel('Date')
```

```
plt.ylabel('Electric Consumption')
plt.legend()
plt.show()
```

	Feature	Importance
5	Factor_F	0.569112
6	Hour	0.183055
1	Factor_B	0.116564
2	Factor_C	0.053908
0	Factor_A	0.020362
8	Month	0.019866
9	Day	0.011156
3	Factor_D	0.008965
4	Factor_E	0.008603
7	DayOfWeek	0.008410





```
[16]: # Compare results
df = pd.DataFrame(results_dict).T
df = df[['cv_rmse', 'cv_rmse_std', 'rmse', 'best_params']]
print(df.to_string())
```

	cv_rmse	cv_rmse_std	rmse	
best_params				
Linear Regression	4.859008	0.664658	4.509047	
Default				
Polynomial Regression	4.190588	1.0516	2.397079	
{'polynomialfeatures__degree': 3, 'polynomialfeatures__include_bias': False}				
Random Forest	2.583831	0.54346	1.633207	
{'max_depth': None, 'min_samples_split': 2, 'n_estimators': 200}				
XGBoost	2.776065	0.557819	1.523946	{'learning_rate': 0.01,
'max_depth': 8, 'n_estimators': 1000, 'predictor': 'cpu_predictor', 'subsample':				
0.8, 'tree_method': 'hist'}				